



POLITECNICO
MILANO 1863

SCUOLA DI INGEGNERIA INDUSTRIALE
E DELL'INFORMAZIONE

NETWORK COMPUTING
Computer Science and Engineering
Ingegneria Informatica

Final Project 2024-25

Author:

Shreesh Kumar Jha

Student ID:

11022306

Advisor:

Prof. Gianni Antichi & Prof. Sebastiano Milano

Academic Year:

2024-25

Contents

Contents	i
1 TCP Connection Tracker Analysis and Performance Report	1
1.1 Source Code Analysis	1
1.2 Critical Bugs Identified	2
1.2.1 Faulty TCP Flag Logic	2
1.2.2 No RST/FIN Handling	2
1.2.3 No Garbage Collection	2
1.2.4 UDP Not Supported	2
1.3 Corner Cases	2
1.4 Implemented Fixes	2
1.4.1 TCP Flag Detection	2
1.4.2 RST & FIN Handling	3
1.4.3 TTL-Based Garbage Collection	3
1.4.4 UDP Flow Tracking	4
1.5 Performance Metrics	4
1.6 Conclusion	4

1 — TCP Connection Tracker Analysis and Performance Report

1.1. Source Code Analysis

In this report and [github](#) repository, I analysed, tested, and debugged the simplified eBPF/XDP TCP connection tracker provided for this project. Originally a basic prototype that evolved into a high-performance system sustaining 3.57 Gbps TCP throughput, I focused on:

TCP State Machine & Handshake Logic: Identified and corrected cases where SYN/SYN-ACK/ACK transitions were mishandled—preventing half-open or simultaneous-open connections from slipping through.

RST/FIN Teardown Handling: Found and fixed missing or incorrect state updates on connection close/reset, ensuring torn-down connections no longer linger in the tracking map.

TTL-Based Garbage Collection: Verified that stale entries are purged reliably, uncovering edge cases where inactive flows persisted past their timeout and tightening expiry checks accordingly.

UDP Flow Tracking: Reviewed minimal UDP “flow” state logic, addressing scenarios where packet bursts could overflow tracking buckets or bypass timeout rules.

Based on this analysis, I implemented targeted fixes—correcting state transitions, closing teardown gaps, reinforcing TTL expiry, and hardening UDP timeouts—to ensure accurate state maintenance under multi-gigabit loads.

1.2. Critical Bugs Identified

1.2.1. Faulty TCP Flag Logic

Issue: Incorrect Boolean expressions prevented SYN-ACK/ACK detection. **Impact:** Broken three-way handshake.

1.2.2. No RST/FIN Handling

Issue: No connection cleanup on termination. **Impact:** “Zombie” entries persisted indefinitely.

1.2.3. No Garbage Collection

Issue: Connections never expired. **Impact:** Map exhaustion under churn.

1.2.4. UDP Not Supported

Issue: All UDP packets dropped. **Impact:** Single-protocol limitation.

1.3. Corner Cases

Handshake Progression: SYN_SENT → SYN_RECV → ESTABLISHED transitions failed

Connection Persistence: No cleanup mechanism for completed connections

Protocol Limitations: TCP-only support limiting deployment

1.4. Implemented Fixes

1.4.1. TCP Flag Detection

```
1 // SYN detection
2 if (pkt.flags == TCPhdr_SYN) { /* Insert SYN_SENT */ }
3
4 // SYN+ACK detection
5 if ((pkt.flags & (TCPhdr_SYN | TCPhdr_ACK)) == (TCPhdr_SYN |
    TCPhdr_ACK)) {
```

```
6      /* Transition SYN_SENT -> SYN_RECV */ }
7
8  // ACK detection
9  if (pkt.flags == TCPHDR_ACK) { /* Transition SYN_RECV ->
    ESTABLISHED */ }
```

Listing 1.1: Corrected flag logic

Result: Complete three-way handshake functionality.

1.4.2. RST & FIN Handling

```
1  // RST handling
2  if (pkt.flags & TCPHDR_RST) {
3      bpf_map_delete_elem(&connections, &key);
4      goto PASS;
5  }
6
7  // FIN state transitions
8  if (saved_state == ESTABLISHED && (pkt.flags & TCPHDR_FIN)) {
9      v->state = FIN_WAIT_1;
10 }
11 if (saved_state == LAST_ACK && pkt.flags == TCPHDR_ACK) {
12     bpf_map_delete_elem(&connections, &key);
13 }
```

Listing 1.2: Connection cleanup

Result: Proper connection termination per RFC 793.

1.4.3. TTL-Based Garbage Collection

```
1  // Expire check on lookup
2  if (v && v->ttl < now) {
3      bpf_map_delete_elem(&connections, &key);
4      v = NULL;
5  }
6
7  // Set TTL on creation
8  newEntry.ttl = now + TCP_SYN_SENT_TIMEOUT; // 2 minutes
```

Listing 1.3: Automatic cleanup

Timeouts: TCP SYN (2min), ESTABLISHED (5 days), FIN (2min), UDP (5-10min).

1.4.4. UDP Flow Tracking

```
1 // New UDP flow
2 newEntry.state = NEW;
3 newEntry.ttl = now + UDP_FLOW_TIMEOUT;
4
5 // Bidirectional detection
6 if ( same_dir) {
7     v->state = ESTABLISHED; // Promote to bidirectional
8 }
```

Listing 1.4: UDP support

Result: 817 Mbps UDP throughput at 1 Gbps offered (4% loss).

1.5. Performance Metrics

Performance Metric	Original	Optimized	Improvement
TCP Throughput (Gbps)	0.283	3.57	12.6x
Connection Stability	1 sec	10+ sec	10x
Packet Retransmissions	5 per test	0	100% reduction
UDP Throughput (Mbps)	0	817	New capability
Protocol Support	TCP only	TCP + UDP	2x protocols

Table 1.1: Comprehensive performance comparison between original and optimized implementations

1.6. Conclusion

The analysis identified four critical bugs preventing optimal functionality and successfully implemented comprehensive fixes. Key achievements:

- Performance:** Achieved 3.57 Gbps TCP throughput (12.6x improvement)
- Protocol Support:** Added UDP with 817 Mbps capability
- Resource Management:** TTL-based garbage collection prevents map exhaustion
- Compliance:** RFC 793-compliant TCP state machine implementation